

But it works on my machine?!

ONNX101



Aleksandar Ivanovski
aleksandar@ivanovski.me

Challenges

- Multiple frameworks
 - Tensorflow
 - PyTorch
- Training on different accelerators
- Inference on various hardware
 - T4 GPUs
 - VPUs

ONNX

- Abstraction and intermediary over frameworks and hardware
- Minimizes the number of moving parts
- Started by AWS, Microsoft and Meta
- Defined format for
 - Deep Learning (DL) models
 - The operators used

Design Principles

- Supports DL and traditional ML
- Flexible enough to keep up with rapid advances
- Compact and cross-platform representation for serialization
- Standardized list of well-defined operators

"ONNX is to Deep Learning what Common Language Runtime is for programming languages"

Linux Foundation AI (LFAI)


- New chapter of Linux Foundation (LF) aimed to drive compatibility and interoperability for AI projects
- Part of LFAI since 2019

ONNX Specification

- Deep Learning (ONNX) - the original
- ONNX-ML aimed for traditional ML (scikit)

ONNX format

- Every computational graph is a direct acyclic graph
- Nodes:
 - Input (known shape)
 - Output (known shape) -> input to the next node
- Operator - essence of the node
- Metadata documents the graph
 - How the model was produced
 - On which accelerator
 - etc...



ONNX Data Types

- Tensor types
 - int8,int16,int32,int64
 - uint8,uint16,uint32,uint64
 - float16,float,double
 - bool
 - string
 - Complex64,complex128
- Non-tensor types in ONNX-ML
 - Sequence
 - Map

ONNX Operators

- Operator is defined by <name, domain, version>
- Operator sets:
 - [ai.onnx](https://ai.onnx.org)
 - [ai.onnx.ml](https://ai.onnx.org/ml)

Operator Examples - Relu

Takes one input data (Tensor) and produces one output data (Tensor) where the rectified linear function, $y = \max(0, x)$, is applied to the tensor element wise.

Version: This version of the operator has been available since version 14 of the ai.onnx operator set.

Other versions of this operator: 1, 6, 13

Inputs: X (differentiable): T (*input tensor*)

Outputs: Y (differentiable): T (*output tensor*)

Type Constraints: T: *tensor - float, int32, int8, int16, int64, float16, double, bfloat16*

Operator Examples - SVMRegressor

Support Vector Machine regression prediction and one-class SVM anomaly detection.

Version: This version of the operator has been available since version 1 of the ai.onnx.ml operator set.

Inputs: T: Data to be regressed.

Outputs: Y (tensor): Regression outputs (one score per target per example).

Attributes: coefficients, kernel_params, kernel_type, n_supports, one_class, ...

Type Constraints: T: *tensor - float, double, int64, int32*

Demo 1


- Define a model in PyTorch
- Train the model with PyTorch
- Export to ONNX
- Visualize the graph

Problem Description


- Train a classifier on the MNIST dataset

2d: 28x28

1d: 1x10




Class: 0



Class: 1

Solution: CNN with PyTorch

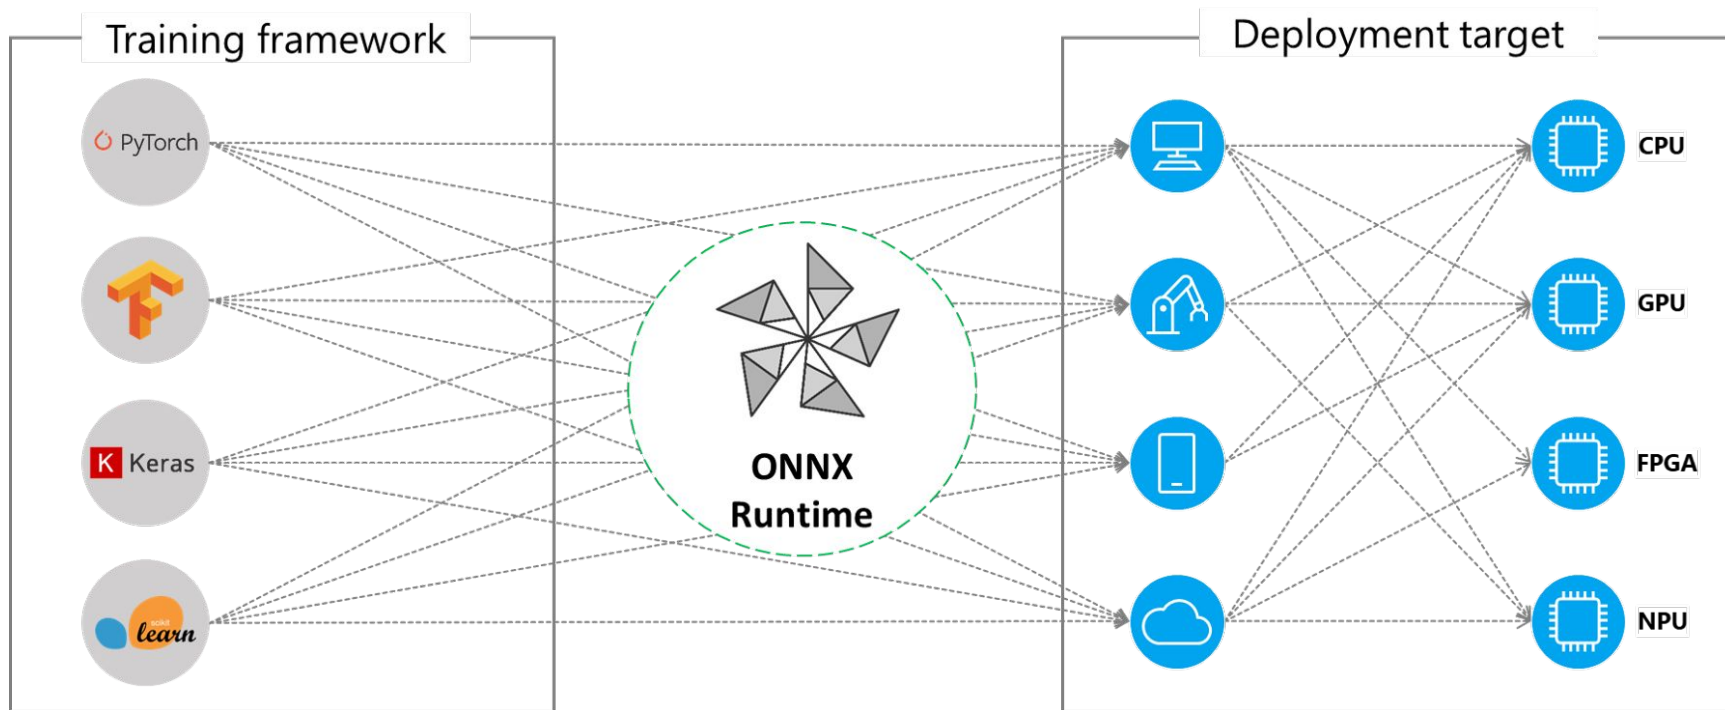
```
class Base42Net(nn.Module):  
    def __init__(self):  
        super(Net, self).__init__()  
        self.conv1 = nn.Conv2d(1, 32, 3, 1)  
        self.conv2 = nn.Conv2d(32, 64, 3, 1)  
        self.dropout1 = nn.Dropout(0.25)  
        self.dropout2 = nn.Dropout(0.5)  
        self.fc1 = nn.Linear(9216, 128)  
        self.fc2 = nn.Linear(128, 10)  
  
    def forward(self, x):  
        x = self.conv1(x)  
        x = F.relu(x)  
        x = self.conv2(x)  
        x = F.relu(x)  
        x = F.max_pool2d(x, 2)  
        x = self.dropout1(x)  
        x = torch.flatten(x, 1)  
        x = self.fc1(x)  
        x = F.relu(x)  
        x = self.dropout2(x)  
        x = self.fc2(x)  
        output = F.log_softmax(x, dim=1)  
        return output
```



ONNX Runtime

- ONNX - file-format and the graph in the file
- ONNX Runtime - how to plug the graph on different hardware
- Created and Maintained by Microsoft

Interoperability



Getting ONNX Models

- [ONNX Model Zoo](#)
- Azure Cognitive Services
- Convert Existing models
- Train from the ground up

Tools

- [Netron](#)
- [VisualDL](#)

Demo 2

- Use the model from Demo 1
- Use onnx-runtime in JS
- Build inference REST API

Express.js inference server

```
import express from 'express'
import ort from 'onnxruntime-node'
import multer from 'multer'
import sharp from 'sharp'

var app = express()
const upload = multer({ storage: multer.memoryStorage() })
const port = 3000

app.post('/classify-digit', upload.single('digit'), async function (req, res) {
  const image = Float32Array.from(await sharp(req.file.buffer).resize(28, 28).grayscale().raw().toBuffer())
  const session = await ort.InferenceSession.create('./models/mnist_cnn.onnx')
  const input = new ort.Tensor('float32', image, [1, 1, 28, 28])

  const output = await session.run({
    "input.1": input
  })
  const predictedNumber = argMax(output['18'].data)
  res.send({
    predictedNumber: predictedNumber
  })
})

app.listen(port, () => {
  console.log(`Starting inference server on ${port}`)
})

const argMax = (array) => {
  return [].reduce.call(array, (m, c, i, arr) => c > arr[m] ? i : m, 0)
}
```

Express.js inference server



~ via 🐉 v3.11.2

```
> curl --request POST \  
  --url http://localhost:3000/classify-digit \  
  --header 'Content-Type: multipart/form-data' \  
  --header 'User-Agent: insomnia/8.3.0' \  
  --form digit=@/examples/3.png  
{"predictedNumber":3}
```

Q&A